

```
In [1]: # Imports
import matplotlib.pyplot as plt
import matplotlib.dates as mdates
import seaborn as sns
import numpy as np
import pandas as pd
from wordcloud import WordCloud
from ucimlrepo import fetch_ucirepo

pd.options.mode.chained_assignment = None
```

```
In [2]: # fetch dataset from the UCI archive website by its ID.
drug_dataset_id = 462
Drugs = fetch_ucirepo(id=drug_dataset_id)
```

```
In [3]: # variable info.
print(Drugs.metadata["additional_info"]["variable_info"])
```

1. drugName (categorical): name of drug
2. condition (categorical): name of condition
3. review (text): patient review
4. rating (numerical): 10 star patient rating
5. date (date): date of review entry
6. usefulCount (numerical): number of users who found review useful

```
In [4]: # Create a total data from the Drugs
data_original = Drugs.data.original

# Checking for missing values in the dataset
missing_values = data_original.isnull().sum()

# Drop the rows with the missing condition value.
data_original = data_original.dropna()

# Basic summary statistics for numerical columns
numerical_summary = data_original.describe()

# Distribution of unique values for categorical columns
drug_count = data_original['drugName'].nunique()
condition_count = data_original['condition'].nunique()

# Summarize text review column (average length of reviews)
data_original['review_length'] = data_original['review'].apply(lambda x: len(x))

average_review_length = data_original['review_length'].mean()

# Display key stats
print(f"Missing Values \n{missing_values}\n")
print(f"Numerical Summary: \n{numerical_summary}\n")
print(f"Number of unique drugs: {drug_count}")
print(f"Number of unique conditions: {condition_count}")
print(f"Average length of review: {round(average_review_length)}")
```

Missing Values

```
id          0
drugName    0
condition   1194
review      0
rating      0
date        0
usefulCount 0
dtype: int64
```

Numerical Summary:

	id	rating	usefulCount
count	213869.000000	213869.000000	213869.000000
mean	116076.924786	6.991149	28.094118
std	67016.705794	3.275792	36.401377
min	0.000000	1.000000	0.000000
25%	58122.000000	5.000000	6.000000
50%	115972.000000	8.000000	16.000000
75%	174018.000000	10.000000	36.000000
max	232291.000000	10.000000	1291.000000

Number of unique drugs: 3667

Number of unique conditions: 916

Average length of review: 85

```
In [5]: # Convert 'date' column to datetime format
data_original['date'] = pd.to_datetime(data_original['date'], format="%d-%b-%Y")

# Extract year from the date
data_original['year'] = data_original['date'].dt.year
# Extract the month from the date
data_original['month'] = data_original['date'].dt.month

# Group data by year and month, and calculate the number of reviews and average rating
data_original['year_month'] = data_original['date'].dt.to_period('M')
```

```
In [6]: # Histogram for numerical variables (rating and usefulCount, assuming they are numerical)
fig, axs = plt.subplots(2, 2, figsize=(12, 10))

# Rating distribution
axs[0, 0].hist(data_original["rating"], bins=10, edgecolor="black", color="lightblue")
axs[0, 0].set_title('Distribution of Ratings')
axs[0, 0].set_xlabel('Rating')
axs[0, 0].set_ylabel('Frequency')

# Rating density plot
sns.kdeplot(data_original['rating'], fill=True, color='blue', ax=axs[0, 1])
axs[0, 1].set_title('Density Plot of Ratings')
axs[0, 1].set_xlabel('Rating')
axs[0, 1].set_ylabel('Density')
plt.grid(True)

# Box plots to show spread and outliers
sns.boxplot(x=data_original["rating"], color="lightblue", ax=axs[1, 0])
axs[1, 0].set_title("Box Plot of Ratings")

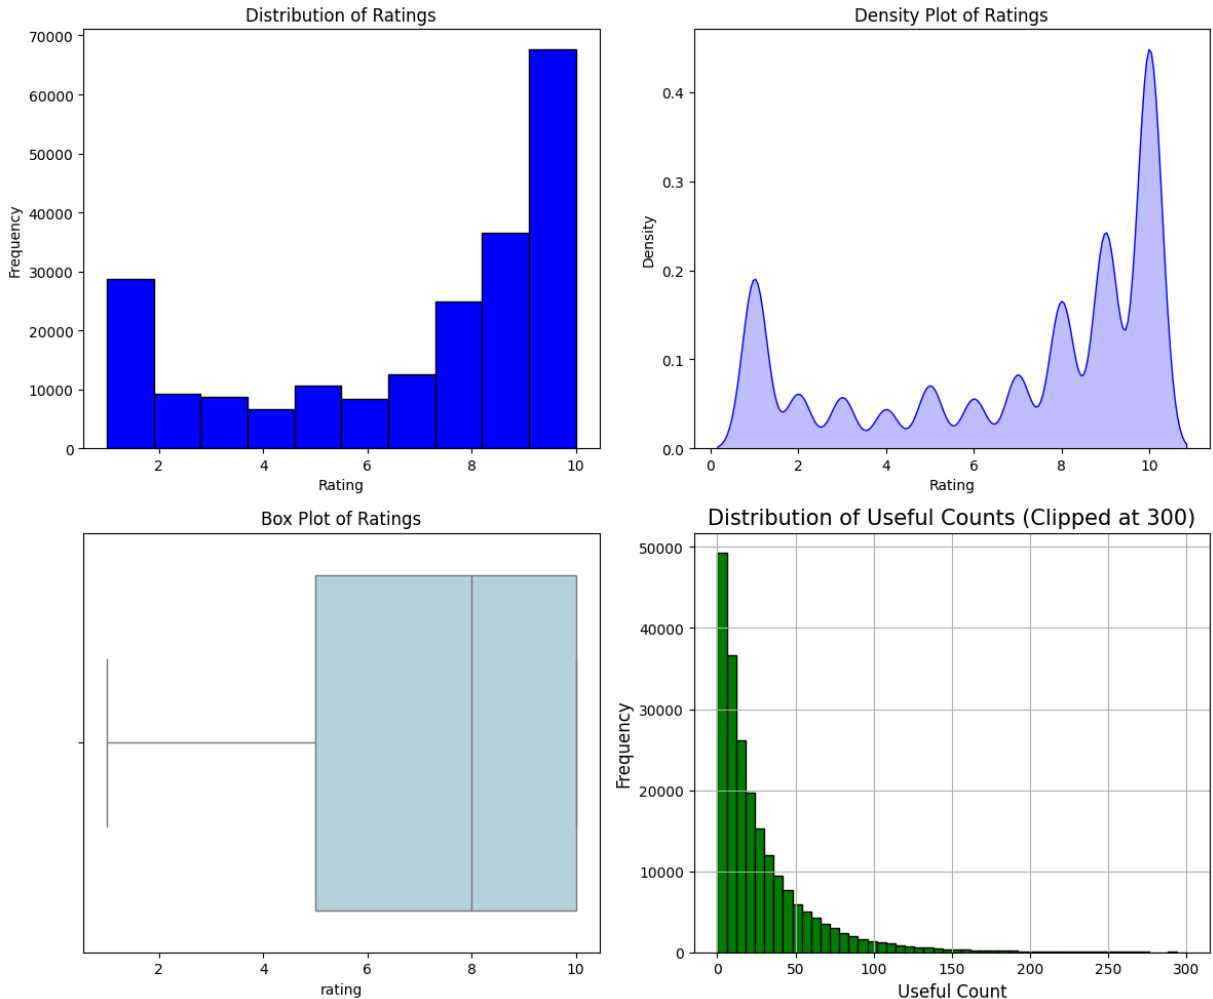
# Distribution of Useful counts plot
axs[1, 1].hist(data_original['usefulCount'], bins=50, color='green', edgecolor='black')
```

```

axs[1, 1].set_title('Distribution of Useful Counts (Clipped at 300)', fontsize=12)
axs[1, 1].set_xlabel('Useful Count', fontsize=12)
axs[1, 1].set_ylabel('Frequency', fontsize=12)
axs[1, 1].grid(True)

plt.tight_layout()
plt.show()

```



```

In [7]: # Calculate the average rating for each condition
avg_rating_per_condition = data_original.groupby("condition")["rating"].mean()
condition_review_counts = data_original.groupby("condition").size() # Count of reviews per condition

conditions_at_least_1000_reviews = condition_review_counts[condition_review_counts >= 1000]
filtered_avg_ratings = avg_rating_per_condition.loc[conditions_at_least_1000_reviews.index]

# Sort the conditions by average rating
sorted_ratings = filtered_avg_ratings.sort_values(ascending=False)

# Visualize the top 10 highest and lowest rated conditions with at least 1000 reviews
plt.figure(figsize=(12, 8))

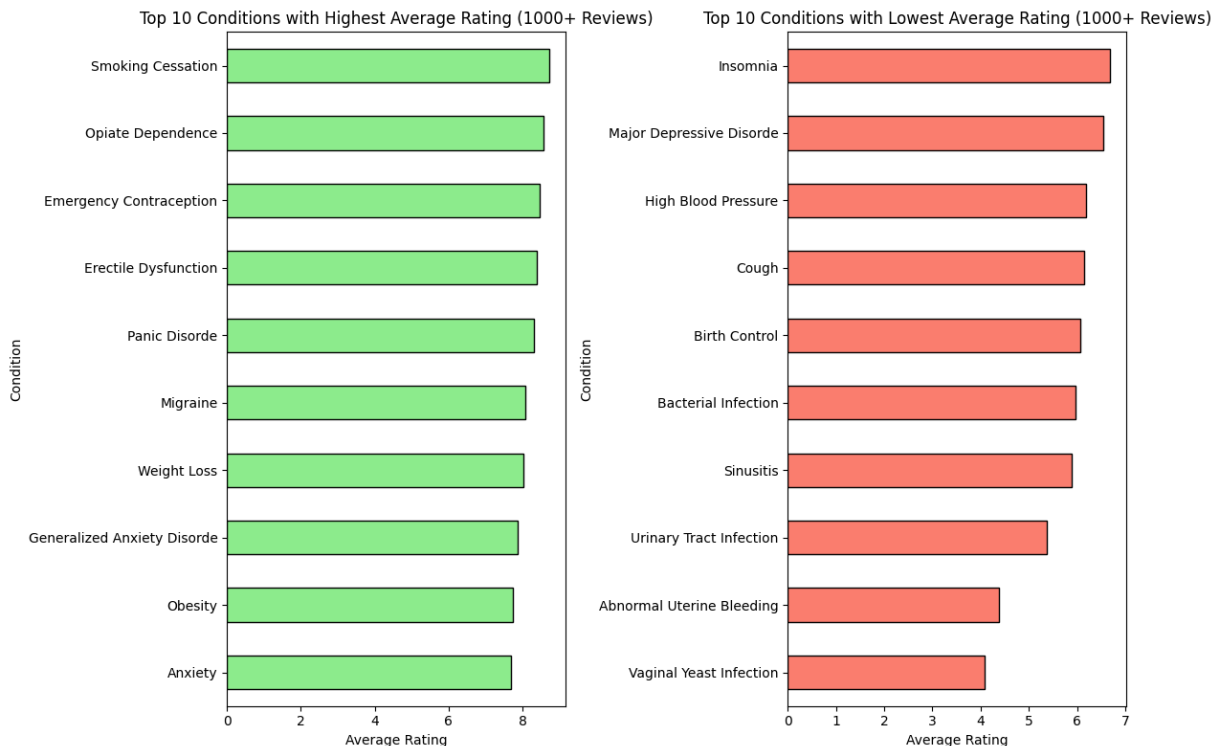
# Plot the top 10 highest rated conditions
plt.subplot(1, 2, 1)
sorted_ratings.tail(10).plot(kind="barh", color="lightgreen", edgecolor="black")
plt.title("Top 10 Conditions with Highest Average Rating (1000+ Reviews)")

```

```
plt.xlabel("Average Rating")
plt.ylabel("Condition")

# Plot the top 10 lowest rated conditions
plt.subplot(1, 2, 2)
sorted_ratings.head(10).plot(kind="barh", color="salmon", edgecolor="black")
plt.title("Top 10 Conditions with Lowest Average Rating (1000+ Reviews)")
plt.xlabel("Average Rating")
plt.ylabel("Condition")

# Display the plots
plt.tight_layout()
plt.show()
```



```
In [8]: # string arr of all of the months in a year, used for the x ticks
months = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
#Including the effect of time aka including the date information
avg_rating_per_drug = data_original.groupby("drugName")["rating"].mean()
drug_review_counts = data_original.groupby("drugName").size()

drugs_at_least_1000_reviews = drug_review_counts[drug_review_counts >= 1000]
filtered_avg_drug_ratings = avg_rating_per_drug.loc[drugs_at_least_1000_reviews]

# Sort the conditions by average rating
sorted_drug_ratings = filtered_avg_drug_ratings.sort_values()

# Identify the top 10 most used drugs (based on number of reviews)
top_10_drugs = sorted_drug_ratings.tail(10).index

top_10_conditions = sorted_ratings.tail(10).index

# Group by drug and year/month, then calculate the average rating
ratings_by_year_top_drugs = data_original[data_original['drugName'].isin(top_10_drugs)]
```

```

ratings_by_month_top_drugs = data_original[data_original['drugName'].isin(top_10_drugs)]

# Group by condition and year/month, then calculate the average rating
ratings_by_year_top_conditions = data_original[data_original['condition'].isin(top_10_conditions)]
ratings_by_month_top_conditions = data_original[data_original['condition'].isin(top_10_conditions)]

# Create a 2x2 grid of subplots
fig, axs = plt.subplots(2, 2, figsize=(12, 10))
fig.suptitle('Drugs/Conditions With At Least 1000 Reviews', fontsize=30)

# Plotting for each drug and condition
for i in range(len(top_10_drugs)):
    drug = top_10_drugs[i]
    sns.lineplot(x='year', y='rating', data=ratings_by_year_top_drugs[drug])
    sns.lineplot(x='month', y='rating', data=ratings_by_month_top_drugs[drug])

    condition = top_10_conditions[i]
    sns.lineplot(x='year', y='rating', data=ratings_by_year_top_conditions[condition])
    sns.lineplot(x='month', y='rating', data=ratings_by_month_top_conditions[condition])

axs[0, 0].set_title('Average Rating Over Time for Top 10 Most Reviewed Drugs')
axs[0, 0].set_xlabel('Year')
axs[0, 0].set_ylabel('Average Rating')

axs[0, 1].set_title('Average Rating Over Time for Top 10 Most Reviewed Drugs')
axs[0, 1].set_xlabel('Month')
axs[0, 1].set_ylabel('Average Rating')
axs[0, 1].set_xticks(range(1, 13)) # Set ticks for each month (1 to 12)
axs[0, 1].set_xticklabels(months)

axs[1, 0].set_title('Average Rating Over Time for Top 10 Most Reviewed Conditions')
axs[1, 0].set_xlabel('Year')
axs[1, 0].set_ylabel('Average Rating')

axs[1, 1].set_title('Average Rating Over Time for Top 10 Most Reviewed Conditions')
axs[1, 1].set_xlabel('Month')
axs[1, 1].set_ylabel('Average Rating')
axs[1, 1].set_xticks(range(1, 13)) # Set ticks for each month (1 to 12)
axs[1, 1].set_xticklabels(months)

# Collect handles and labels from one of the subplots
handles_drugs, labels_drugs = axs[0, 0].get_legend_handles_labels()
handles_condition, labels_condition = axs[1, 0].get_legend_handles_labels()

# Remove the legends as it will be displayed on the side
axs[0, 0].legend().remove()
axs[0, 1].legend().remove()
axs[1, 0].legend().remove()
axs[1, 1].legend().remove()

# Set a common legend outside the plot (to the right)
fig.legend(handles_drugs, labels_drugs, loc='center left', bbox_to_anchor=(1, 0.5))
fig.legend(handles_condition, labels_condition, loc='center left', bbox_to_anchor=(1, 0.5))


# Rotate the x-ticks to make them readable

```

```
plt.setp(axes[0, 1].get_xticklabels(), rotation=45, ha="right")
plt.setp(axes[1, 1].get_xticklabels(), rotation=45, ha="right")

plt.tight_layout()
plt.show()
```

Drugs/Conditions With At Least 1000 Reviews



```
In [9]: # Aggregating data by month (number of reviews and average rating)
temporal_data = data_original.groupby('year_month').agg(
    num_reviews=('review', 'count'),
    avg_rating=('rating', 'mean')
).reset_index()

# Convert year_month back to datetime for proper plotting
temporal_data['year_month'] = temporal_data['year_month'].dt.to_timestamp()

# Plotting temporal trends
fig, ax1 = plt.subplots(figsize=(12, 6))


# Plot number of reviews over time
ax1.plot(temporal_data['year_month'], temporal_data['num_reviews'], color='b')
ax1.set_xlabel('Date', fontsize=12)
ax1.set_ylabel('Number of Reviews', fontsize=12, color='blue')
ax1.tick_params(axis='y', labelcolor='blue')

# Create a second y-axis for the average rating
ax2 = ax1.twinx()
ax2.plot(temporal_data['year_month'], temporal_data['avg_rating'], color='g')
ax2.set_ylabel('Average Rating', fontsize=12, color='green')
ax2.tick_params(axis='y', labelcolor='green')

plt.title('Temporal Trend of Reviews and Average Rating', fontsize=15)
plt.grid(True)
```

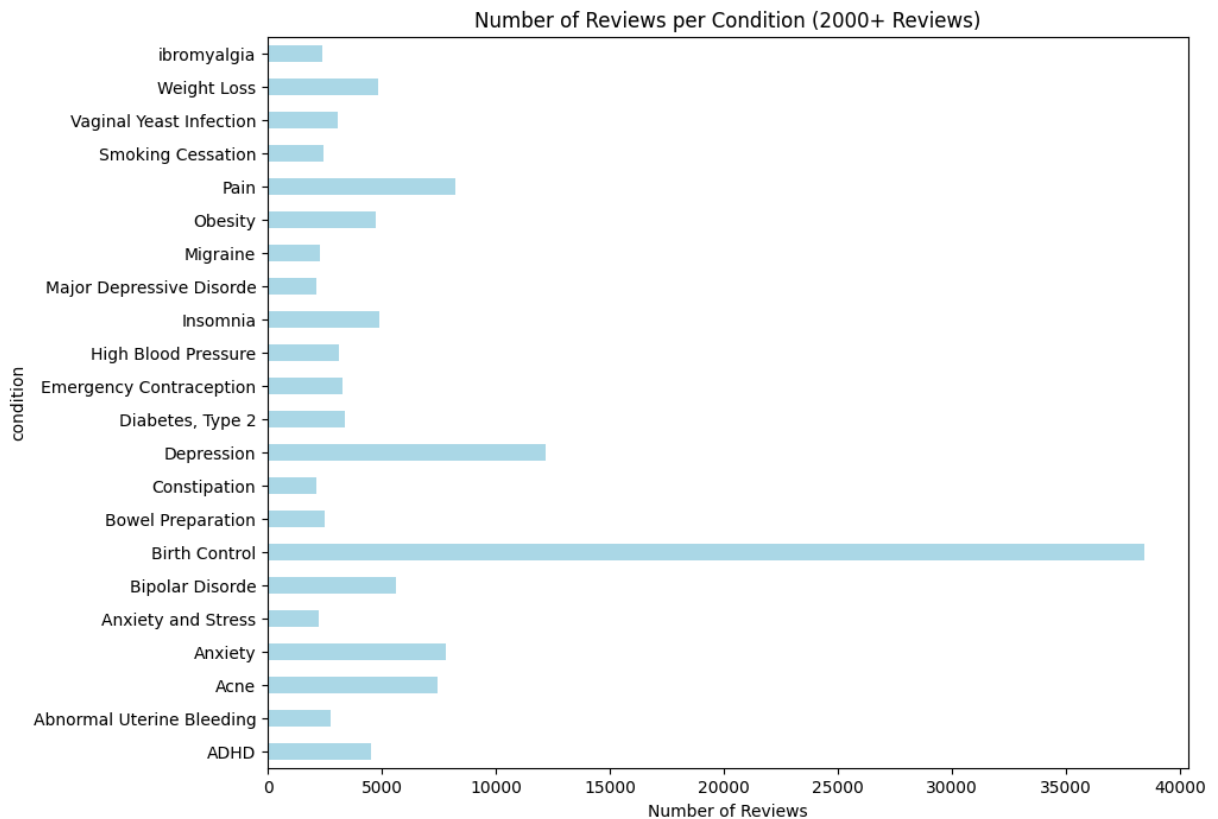
```
plt.tight_layout()
```

```
plt.show()
```



```
In [10]: conditions_large_reviews = condition_review_counts[condition_review_counts > 2000]


plt.figure(figsize=(12, 5))
conditions_large_reviews.plot(kind="barh", figsize=(10, 8), color="lightblue")
plt.title("Number of Reviews per Condition (2000+ Reviews)")
plt.xlabel("Number of Reviews")
plt.show()
```



```
In [11]: # Filter conditions with at least 2000 reviews
conditions_at_least_2000_reviews = condition_review_counts[condition_review_
filtered_avg_ratings_2000 = avg_rating_per_condition.loc[conditions_at_least

plt.figure(figsize=(10, 6))
filtered_avg_ratings_2000.sort_values().plot(kind="barh", color="skyblue", e
plt.title("Average Rating for Conditions with At Least 2000 Reviews")
plt.xlabel("Average Rating")
plt.ylabel("Condition")
plt.tight_layout()

# Show the plot
plt.show()
```



```
In [12]: adhd = data_original[data_original["condition"] == "ADHD"]

plt.figure(figsize=(12, 5))

# Histogram for "rating"
plt.subplot(1, 2, 1)
plt.hist(adhd["rating"], bins=10, color="blue", edgecolor="black")
plt.title("Distribution of Ratings")
plt.xlabel("Rating")
plt.ylabel("Frequency")



# Generate a string of all drug names repeated based on their counts
drug_counts = adhd["drugName"].head(30).value_counts()
drug_string = " ".join([f"{drug} " * count for drug, count in drug_counts.items()])

# Create the word cloud
wordcloud = WordCloud(width=1000, height=800, background_color="white").gene
```



```
# Plot the word cloud
plt.subplot(1, 2, 2)
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off") # Turn off the axis
plt.title("Word Cloud of Drugs by Number of Reviews")
plt.show()

# Display the plots
plt.tight_layout()
plt.show()
```

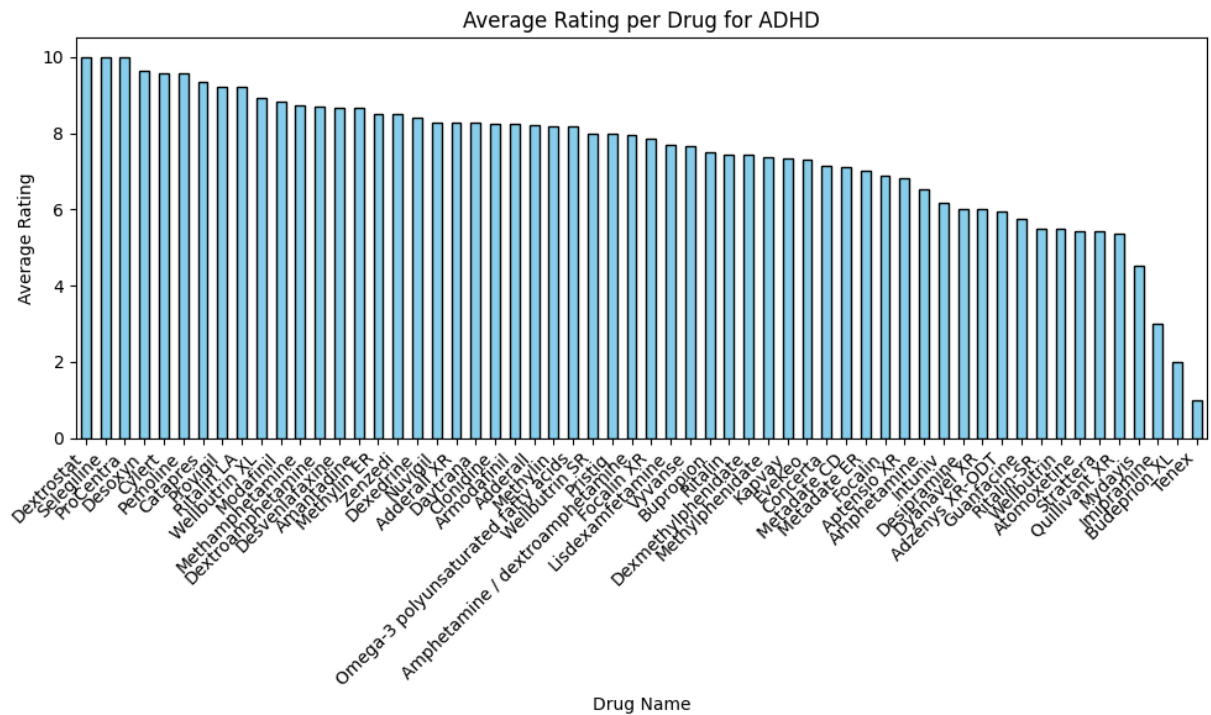


<Figure size 640x480 with 0 Axes>

```
In [13]: # Calculate the average rating for each drug
avg_rating_per_drug = adhd.groupby("drugName")["rating"].mean()

# Plotting the average rating per drug
plt.figure(figsize=(10, 6))
avg_rating_per_drug.sort_values(ascending=False).plot(kind="bar", color="sky")
plt.title("Average Rating per Drug for ADHD")
plt.xlabel("Drug Name")
plt.ylabel("Average Rating")
plt.xticks(rotation=45, ha="right")
plt.tight_layout()

# Show the plot
plt.show()
```



```
In [14]: # Import local module helper function that checks if side effects are mentioned
from utils.side_effects import check_side_effects


# Apply the function to the review text and create the 'side_effects' column
data_original['side_effects'] = data_original['review'].apply(check_side_effects)

# Check the distribution of the new target variable
side_effect_counts = data_original['side_effects'].value_counts()
#print(f"Side Effect Target Distribution:\n{side_effect_counts}")

# Group by drug name and sum the 'side_effects' column to count how many reviews mention side effects
side_effects_by_drug = data_original.groupby('drugName')['side_effects'].sum()

# Sort the results to find the drugs with the most side effects
side_effects_by_drug_sorted = side_effects_by_drug.sort_values(by='side_effects', ascending=False)

# Plot the top 10 drugs with the most side effects
plt.figure(figsize=(10, 6))
sns.barplot(x='side_effects', y='drugName', data=side_effects_by_drug_sorted)
plt.title('Top 10 Drugs with the Most Side Effects')
plt.xlabel('Number of Reviews Mentioning Side Effects')
plt.ylabel('Drug Name')
plt.show()
```



```
In [15]: from nltk.sentiment.vader import SentimentIntensityAnalyzer
from sklearn.model_selection import train_test_split
sid = SentimentIntensityAnalyzer()

# Split the data into 75% training and 25% testing
train_df, test_df = train_test_split(data_original, test_size=0.25, random_s
# Preprocess text for sentiment analysis (get sentiment score from review te
data_original['sentiment'] = data_original['review'].apply(lambda x: sid.pol

# Separate the data into two groups: negative (sentiment < 0) and neutral/po
negative = data_original[data_original['sentiment'] < 0]
neutral_positive = data_original[data_original['sentiment'] >= 0]


avg_rating_negative = negative['rating'].mean()
avg_rating_neutral_positive = neutral_positive['rating'].mean()

# Print the results
print(f"Average rating for negative sentiment (sentiment < 0): {avg_rating_r
print(f"Average rating for neutral/positive sentiment (sentiment >= 0): {avg

# plot the distribution of ratings for each sentiment group
plt.figure(figsize=(10, 6))
sns.histplot(negative['rating'], color='red', label='Negative Sentiment (< 0)')
sns.histplot(neutral_positive['rating'], color='green', label='Neutral/Positive Sentiment (>= 0)')
plt.title('Distribution of Ratings by Sentiment Groups')
plt.xlabel('Rating')
plt.ylabel('Density')
plt.legend()
plt.show()
```

Average rating for negative sentiment (sentiment < 0): 6.08

Average rating for neutral/positive sentiment (sentiment >= 0): 7.93



```
In [16]: from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Define effectiveness for training data, here we assumed the ratings > 5 me
# which can be adjusted with the previous section of sentiment + rating
train_df['effective'] = train_df.apply(
    lambda row: 1 if row['rating'] > 5 else 0, axis=1
)

# Define effectiveness for test data in the same way
test_df['effective'] = test_df.apply(
    lambda row: 1 if row['rating'] > 5 else 0, axis=1
)

# Handle missing data
train_df.dropna(subset=['drugName', 'condition', 'rating'], inplace=True)
test_df.dropna(subset=['drugName', 'condition', 'rating'], inplace=True)

# Encode categorical variables (Drug name, Condition)
le_drug = LabelEncoder()
le_condition = LabelEncoder()

train_df['drug_encoded'] = le_drug.fit_transform(train_df['drugName'])
train_df['condition_encoded'] = le_condition.fit_transform(train_df['condition'])

# Define a safe transformation function to handle unseen labels
def safe_transform(encoder, data, default_value=-1):
    return np.array([default_value if label not in encoder.classes_ else encoder.transform([label])[0] for label in data])

# Apply safe transformation on the test data to handle unseen drugs and conditions
test_df['drug_encoded'] = safe_transform(le_drug, test_df['drugName'])
```

```

test_df['condition_encoded'] = safe_transform(le_condition, test_df['condition'])

# Define the input features (drug name and condition) and the target variable
X_train = train_df[['drug_encoded', 'condition_encoded']]
y_train = train_df['effective']

X_test = test_df[['drug_encoded', 'condition_encoded']]
y_test = test_df['effective']

# Train a Random Forest classifier using drug name and condition as input features
clf = RandomForestClassifier(n_estimators=100, random_state=42)
clf.fit(X_train, y_train)

# Make predictions on the test set using drug name and condition
y_pred = clf.predict(X_test)

# Evaluate the model
print(f"Train Accuracy: {accuracy_score(y_train, clf.predict(X_train))}")
print(f"Test Accuracy: {accuracy_score(y_test, y_pred)}")
print(classification_report(y_test, y_pred))

# Create a comparison table of actual vs predicted values
comparison_df = pd.DataFrame({
    'Drug Name': test_df['drugName'],
    'Condition': test_df['condition'],
    'Actual Effectiveness': y_test,
    'Predicted Effectiveness': y_pred
})

print(comparison_df.head())

# Plot the predicted effectiveness for a few drug-condition pairs
plt.figure(figsize=(10, 6))
sns.countplot(x='Actual Effectiveness', hue='Predicted Effectiveness', data=comparison_df)
plt.title('Comparison of Actual vs Predicted Effectiveness')
plt.show()

```


Train Accuracy: 0.7325016676953385

Test Accuracy: 0.7094710855090895

	precision	recall	f1-score	support
0	0.54	0.19	0.29	16035
1	0.73	0.93	0.82	37433
accuracy			0.71	53468
macro avg	0.64	0.56	0.55	53468
weighted avg	0.67	0.71	0.66	53468

	Drug Name	Condition	Actual Effectiveness \
148619	Epiduo	Acne	1
90835	Levofloxacin	Bronchitis	0
194827	Amlodipine	High Blood Pressure	1
113551	Dapsone	Acne	0
89748	Amitriptyline	Insomnia	1

	Predicted Effectiveness
148619	1
90835	0
194827	0
113551	1
89748	1



```
In [17]: #Some visualizations, we can pick and chose what we want to keep
from sklearn.metrics import confusion_matrix
import seaborn as sns

# Train and Test Accuracy as a bar plot
plt.figure(figsize=(8, 5))
accuracy_values = [accuracy_score(y_train, clf.predict(X_train)), accuracy_s
plt.bar(['Train Accuracy', 'Test Accuracy'], accuracy_values, color=['blue',
```


```

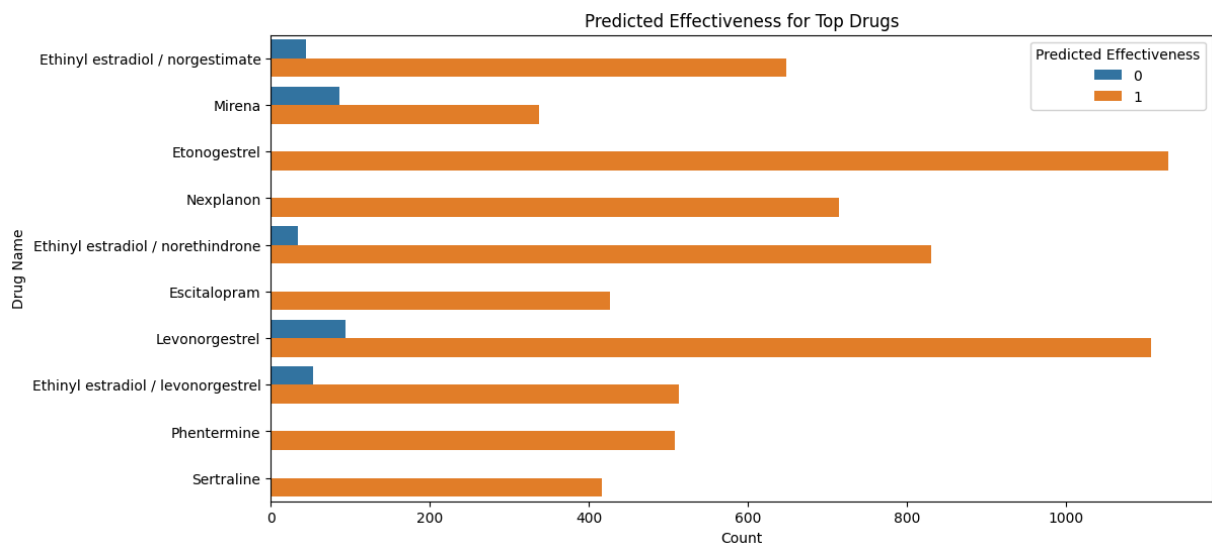

plt.title('Train vs Test Accuracy')
plt.ylabel('Accuracy')
plt.ylim(0, 1)
plt.show()

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False, xticklabels=[
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# Bar plot of Predicted Effectiveness for different drugs
plt.figure(figsize=(12, 6))
top_drugs = comparison_df['Drug Name'].value_counts().index[:10]
sns.countplot(y='Drug Name', hue='Predicted Effectiveness', data=comparison_
plt.title('Predicted Effectiveness for Top Drugs')
plt.xlabel('Count')
plt.ylabel('Drug Name')
plt.legend(title='Predicted Effectiveness', loc='upper right')
plt.show()

```





```
In [18]: # Some additional descriptive analysis that could be useful:
# Select all numerical columns except 'sentiment'
numerical_columns = data_original.select_dtypes(include=['float64', 'int64'])
numerical_columns = numerical_columns.drop('sentiment', errors='ignore') #

# Calculate the correlation matrix for the selected numerical features
correlation_matrix = data_original[numerical_columns].corr()

# Display the correlation matrix
print("Correlation Matrix (without Sentiment):")
```



```

print(correlation_matrix)

# Plot a heatmap of the correlation matrix
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Correlation Matrix of All Numerical Features (without Sentiment)')
plt.show()

# Scatter plot for rating vs usefulCount (helpful votes)
plt.figure(figsize=(8, 6))
sns.scatterplot(x='rating', y='usefulCount', data=data_original)
plt.title('Rating vs Helpful Votes')
plt.xlabel('Rating')
plt.ylabel('Helpful Votes')
plt.show()


```

Correlation Matrix (without Sentiment):

	id	rating	usefulCount	review_length	side_effects
id	1.000000	0.015925	0.017024	0.007668	-0.002079
rating	0.015925	1.000000	0.235121	0.026788	-0.095272
usefulCount	0.017024	0.235121	1.000000	0.019912	-0.009833
review_length	0.007668	0.026788	0.019912	1.000000	0.303126
side_effects	-0.002079	-0.095272	-0.009833	0.303126	1.000000

Correlation Matrix of All Numerical Features (without Sentiment)





```
In [19]: # Top 10 most common drugs
top_drugs = data_original['drugName'].value_counts().head(10)
print("Top 10 Most Reviewed Drugs:")
print(top_drugs)

# Visualize the top 10 most common drugs
plt.figure(figsize=(10, 6))
sns.barplot(x=top_drugs.values, y=top_drugs.index, hue=top_drugs.index, legend=True)
plt.title('Top 10 Most Reviewed Drugs')
plt.xlabel('Number of Reviews')
plt.ylabel('Drug Name')
plt.show()

# Top 10 most common conditions
top_conditions = data_original['condition'].value_counts().head(10)
print("Top 10 Most Reviewed Conditions:")
print(top_conditions)

# Visualize the top 10 conditions
plt.figure(figsize=(10, 6))
sns.barplot(x=top_conditions.values, y=top_conditions.index, hue=top_conditions.index, legend=True)
plt.title('Top 10 Most Reviewed Conditions')
plt.xlabel('Number of Reviews')
plt.ylabel('Condition')
plt.show()

# Number of reviews per drug and condition
```

```

reviews_per_drug_condition = data_original.groupby(['drugName', 'condition'])
top_drug_condition = reviews_per_drug_condition.sort_values(by='num_reviews')

print("Top 10 Drug-Condition Pairs by Number of Reviews:")
print(top_drug_condition)

# Visualize the top drug-condition pairs
plt.figure(figsize=(10, 6))
sns.barplot(x='num_reviews', y='drugName', hue='condition', data=top_drug_cc)
plt.title('Top 10 Drug-Condition Pairs by Number of Reviews')
plt.xlabel('Number of Reviews')
plt.ylabel('Drug Name')
plt.legend(title='Condition', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.show()


# Plot the distribution of sentiment scores
plt.figure(figsize=(10, 6))
sns.histplot(data_original['sentiment'], bins=20, kde=True, color='purple')
plt.title('Sentiment Score Distribution')
plt.xlabel('Sentiment Score')
plt.ylabel('Count')
plt.show()

# Show how many reviews have positive, negative, or neutral sentiment
print("Sentiment Distribution:")
sentiment_labels = ['Negative', 'Neutral', 'Positive']
sentiment_distribution = [
    (data_original['sentiment'] < 0).sum(),
    (data_original['sentiment'] == 0).sum(),
    (data_original['sentiment'] > 0).sum()
]
for label, count in zip(sentiment_labels, sentiment_distribution):
    print(f"{label}: {count} reviews")

```

Top 10 Most Reviewed Drugs:

drugName	
Levonorgestrel	4896
Etonogestrel	4402
Ethinyl estradiol / norethindrone	3619
Nexplanon	2892
Ethinyl estradiol / norgestimate	2682
Ethinyl estradiol / levonorgestrel	2400
Phentermine	2077
Sertraline	1859
Escitalopram	1739
Mirena	1673
Name: count, dtype: int64	




Top 10 Most Reviewed Conditions:

condition


Birth Control	38436
Depression	12164
Pain	8245
Anxiety	7812
Acne	7435
Bipolar Disorde	5604
Insomnia	4904
Weight Loss	4857
Obesity	4757
ADHD	4509

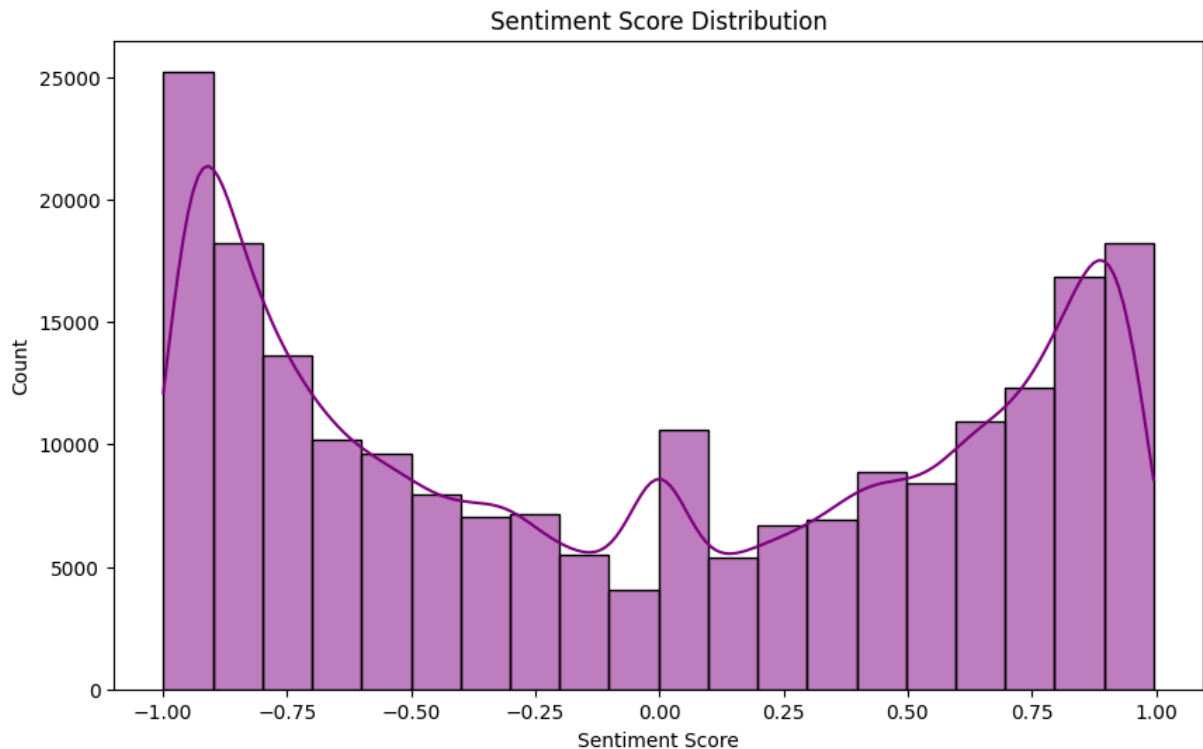
Name: count, dtype: int64



Top 10 Drug-Condition Pairs by Number of Reviews:

	drugName	condition	num_revie
WS			
3307	Etonogestrel	Birth Control	43
94			
3283	Ethinyl estradiol / norethindrone	Birth Control	30
81			
4790	Levonorgestrel	Birth Control	28
84			
6012	Nexplanon	Birth Control	28
83			
3274	Ethinyl estradiol / levonorgestrel	Birth Control	21
07			
3292	Ethinyl estradiol / norgestimate	Birth Control	20
97			
4791	Levonorgestrel	Emergency Contraception	16
51			
6693	Phentermine	Weight Loss	16
50			
4233	Implanon	Birth Control	14
96			
5527	Miconazole	Vaginal Yeast Infection	13
38			





Sentiment Distribution:
 Negative: 108677 reviews
 Neutral: 6728 reviews
 Positive: 98464 reviews

```
In [20]: import ipywidgets as widgets
from IPython.display import display

# Define effectiveness using percentiles of the ratings to add more variance
# Compute percentiles and scale effectiveness
data_original['percentile_effectiveness'] = data_original.groupby('condition')

# Group by condition and drug, then calculate the mean percentile effectiveness
effectiveness_by_condition = data_original.groupby(['condition', 'drugName'])
    average_rating=('rating', 'mean'),
    confidence=('percentile_effectiveness', 'mean') # Now using percentile
).reset_index()

# Function to get the best drug for a given condition
def get_best_drug(condition):
    condition_data = effectiveness_by_condition[condition]

    if condition_data.empty:
        return "No data available", 0

    best_drug = condition_data.sort_values(by='confidence', ascending=False)

    return best_drug['drugName'], best_drug['confidence'] # Return the conf

# Create a dropdown widget for selecting a condition
condition_dropdown = widgets.Dropdown(
    options=data_original['condition'].unique(),
    description='Condition:',
```

```

        disabled=False
    )

    # Create a label widget to display the best drug and confidence
    result_label = widgets.Label(value="Select a condition to see the best drug")

    # Function to update the result when a condition is selected
    def on_condition_change(change):
        if change['type'] == 'change' and change['name'] == 'value':
            selected_condition = change['new']
            best_drug, confidence = get_best_drug(selected_condition)
            result_label.value = f"Best Drug: {best_drug}\nConfidence: {confidence}"

    # Attach the function to the dropdown widget
    condition_dropdown.observe(on_condition_change)

    # Display the dropdown and result label in the notebook
    display(condition_dropdown)
    display(result_label)

```

```

Dropdown(description='Condition:', options=('Left Ventricular Dysfunction',
'ADHD', 'Birth Control', 'Opiate D...
Label(value='Select a condition to see the best drug')

```

This notebook was converted with convert.ploomber.io